

# Intro to Deep Learning

Big Data and Machine Learning en el Mercado Inmobiliario  
Educación Continua

Ignacio Sarmiento-Barbieri

Universidad de los Andes

November 3, 2022

# Deep Learning: Intro

- ▶ Neural networks are simple models.
- ▶ Their strength lays in their simplicity
- ▶ Neural networks combine inputs that are passed through nonlinear activation functions called nodes (or, in reference to the human brain, neurons), to approximate  $f^*(x)$

# Deep Learning: Intro

- ▶ Let's start with a familiar and simple model, the linear model

$$y = f(X) + u \tag{1}$$


$$y = \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + u$$

# Deep Learning: Intro

- Let's start with a familiar and simple model, the linear model

$$y = f(X) + u \quad (1)$$

$$y = \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + u$$



# Single Layer Neural Networks

- ▶ Linear Models may be too simple, and miss the nonlinearities that best approximate  $f^*(x)$
- ▶ We can overcome these limitations of linear models and handle a more general class of functions by incorporating one or more hidden layers.
- ▶ Neural Networks are also called deep feedforward networks, feedforward neural networks, or multilayer perceptrons (MLPs), and are the quintessential deep learning models

# Single Layer Neural Networks

- ▶ A neural network takes an input vector of  $p$  variables

$$X = (X_1, X_2, \dots, X_p) \quad (2)$$

- ▶ and builds a nonlinear function  $f(X)$  to predict the response  $y$ .

$$y = f(X) + u \quad (3)$$

- ▶ What distinguishes neural networks from previous methods is the particular structure of the model.

# Single Layer Neural Networks

- ▶ A NN model has the form

$$f(X) = \beta_0 + \sum_{k=1}^K \beta_k A_k \quad (4)$$

$$= \beta_0 + \sum_{k=1}^K \beta_k h_k(X) \quad (5)$$


$$= \beta_0 + \sum_{k=1}^K \beta_k g \left( w_{k0} + \sum_{j=1}^p w_{kj} X_j \right) \quad (6)$$

- ▶ where  $g(\cdot)$  is a activation function specified in advance
- ▶ where the nonlinearity of  $g(\cdot)$  is essential

# Single Layer Neural Networks

► Let's consider a very simple example with

- $p = 2, X = (X_1, X_2)$
- $K = 2, h_1(X)$  and  $h_2(X)$
- $g(z) = z^2$





# Single Layer Neural Networks

- ▶ Let's consider a very simple example with
  - ▶  $p = 2, X = (X_1, X_2)$
  - ▶  $K = 2, h_1(X)$  and  $h_2(X)$
  - ▶  $g(z) = z^2$
- ▶ Then

$$f(X) = \beta_0 + \sum_{k=1}^2 \beta_k A_k \quad (7)$$

$$= \beta_0 + \sum_{k=1}^2 \beta_k h_k(X) \quad (8)$$

$$= \beta_0 + \sum_{k=1}^2 \beta_k g \left( w_{k0} + \sum_{j=1}^p w_{kj} X_j \right) \quad (9)$$

# Single Layer Neural Networks

$$f(X) = \beta_0 + \sum_{k=1}^2 \beta_k g \left( w_{k0} + \sum_{j=1}^2 w_{kj} X_j \right) \quad (10)$$

► We specify the parameters as

$$\begin{array}{lll} \beta_0 = 0 & \beta_1 = \frac{1}{4} & \beta_2 = -\frac{1}{4} \\ w_{10} = 0 & w_{11} = 1 & w_{12} = 1 \\ w_{20} = 0 & w_{21} = 1 & w_{22} = -1 \end{array}$$

# Single Layer Neural Networks

► Then

$$h_1(X) = (0 + X_1 + X_2)^2 \quad (11)$$

$$h_2(X) = (0 + X_1 - X_2)^2 \quad (12)$$

► and plugging in

$$f(X) = 0 + \frac{1}{4} (0 + X_1 + X_2)^2 - \frac{1}{4} (0 + X_1 - X_2)^2 \quad (13)$$

$$= \frac{1}{4} \left( (X_1 + X_2)^2 - (X_1 - X_2)^2 \right) \quad (14)$$

$$= X_1 X_2 \quad (15)$$

## Worked Example: The "Exclusive OR (XOR)" Function

- ▶ The exclusive disjunction of a pair of propositions,  $(p, q)$ , is supposed to mean that  $p$  is true or  $q$  is true, but not both
- ▶ It's truth table is:

$q$	$p$	$q \vee p$
0	0	0
0	1	1
1	0	1
1	1	0

- ▶ When exactly one of these binary values is equal to 1, the XOR function returns 1. Otherwise, it returns 0

## Worked Example: The "Exclusive OR (XOR)" Function

- ▶ Let's use a linear model

$$y = X\beta + \iota\alpha \quad (16)$$


$$y = \begin{pmatrix} 0 \\ 1 \\ 1 \\ 0 \end{pmatrix} \quad X = \begin{pmatrix} 0 & 0 \\ 0 & 1 \\ 1 & 0 \\ 1 & 1 \end{pmatrix} \quad \iota = \begin{pmatrix} 1 \\ 1 \\ 1 \\ 1 \end{pmatrix} \quad (17)$$

- ▶ Solution  $\alpha = \frac{1}{2}$   $\beta = \begin{pmatrix} 0 \\ 0 \end{pmatrix}$

- ▶ Prediction  $\hat{y} = \begin{pmatrix} \frac{1}{2} \\ \frac{1}{2} \\ \frac{1}{2} \\ \frac{1}{2} \end{pmatrix}$

# Worked Example: The "Exclusive OR (XOR)" Function

- ▶ Let's use Single Layer NN containing two hidden units



## Worked Example: The “Exclusive OR (XOR)” Function

- ▶ Which activation functions ( $A_k$ ) should we choose?
  - ▶ Clearly **not** linear, otherwise it would defeat the entire purpose
  - ▶ We are going to use the rectified linear unit or ReLU (it is usually the default recommendation, there are many others (more on this later))
  - ▶ ReLU is defined as  $g(z) = \max\{0, z\}$
- ▶ For the output layer? For this example, a linear model will suffice

$$f(X) = \beta_0 + w_k A_k \quad (18)$$

- ▶ The final model is then

$$f(x, W, C, w, b) = \max\{0, XW + c\} w + b \quad (19)$$

## Worked Example: The "Exclusive OR (XOR)" Function

- Suppose this is the solution to the XOR problem

$$f(x) = \max\{0, XW + c\} w + b$$

$$W = \begin{pmatrix} 1 & 1 \\ 1 & 1 \end{pmatrix}$$

$$c = \begin{pmatrix} 0 & -1 \\ 0 & -1 \\ 0 & -1 \\ 0 & -1 \end{pmatrix}$$

$$w = (1 \quad -2)$$

$$b = 0$$



## Worked Example: The "Exclusive OR (XOR)" Function

- Lets work out the example step by step

$$f(x) = \max\{0, XW + c\} w + b \quad (20)$$

$$XW = \begin{pmatrix} 0 & 0 \\ 0 & 1 \\ 1 & 0 \\ 1 & 1 \end{pmatrix} \begin{pmatrix} 1 & 1 \\ 1 & 1 \end{pmatrix} = \begin{pmatrix} 0 & 0 \\ 1 & 1 \\ 1 & 1 \\ 2 & 2 \end{pmatrix}$$

$$XW + c = \begin{pmatrix} 0 & -1 \\ 1 & 0 \\ 1 & 0 \\ 2 & 1 \end{pmatrix}$$

$$\max\{0, XW + c\} = \begin{pmatrix} 0 & 0 \\ 1 & 0 \\ 1 & 0 \\ 2 & 1 \end{pmatrix}$$

## Worked Example: The "Exclusive OR (XOR)" Function

$$\hat{y} = \max\{0, XW + c\} w + b = \begin{pmatrix} 0 & 0 \\ 1 & 0 \\ 1 & 0 \\ 2 & 1 \end{pmatrix} \begin{pmatrix} 1 \\ -2 \end{pmatrix} = \begin{pmatrix} 0 \\ 1 \\ 1 \\ 0 \end{pmatrix}$$

- The neural network has obtained the correct answer for every data point

## Worked Example: The “Exclusive OR (XOR)” Function

- ▶ In this example, we simply specified the solution, then showed that it obtained zero error.
- ▶ In a real situation, obviously we can't guess the solution
- ▶ What we do is gradient based optimization
- ▶ Remember that the convergence point of gradient descent depends on the initial values of the parameters and step size.
- ▶ In practice, gradient descent would usually not find clean, easily understood, integer-valued solutions like we did here.

# NN Minimalist Theory

- ▶ Why not a linear activation functions?
- ▶ Let's go back to our example
  - ▶  $p = 2, X = (X_1, X_2)$
  - ▶  $K = 2, h_1(X)$  and  $h_2(X)$
  - ▶ Now  $g(z) = z$
- ▶ Then

$$f(X) = \beta_0 + \sum_{k=1}^2 \beta_k A_k \quad (21)$$

$$= \beta_0 + \sum_{k=1}^2 \beta_k h_k(X) \quad (22)$$

$$= \beta_0 + \sum_{k=1}^2 \beta_k g \left( w_{k0} + \sum_{j=1}^p w_{kj} X_j \right) \quad (23)$$

# NN Minimalist Theory

Why not a linear activation functions?

- ▶ Since  $g(z) = z$  we get

$$f(X) = \beta_0 + \sum_{k=1}^2 \beta_k \left( w_{k0} + \sum_{j=1}^2 w_{kj} X_j \right) \quad (24)$$

- ▶ Replacing

$$f(X) = \beta_0 + \beta_1 (w_{10} + w_{11}X_1 + w_{12}X_2) + \beta_2 (w_{20} + w_{21}X_1 + w_{22}X_2) \quad (25)$$

- ▶ then

$$f(X) = \theta_0 + \theta_1 X_1 + \theta_2 X_2 \quad (26)$$

# Activation Functions

- ▶ The gain comes from using nonlinear activation function  $f$
- ▶ Note that, with nonlinear activation functions in place, it is no longer possible to collapse our NN into a linear model.
- ▶ Activation functions are fundamental to deep learning, let us briefly survey some common activation functions.
- ▶ In practice we would not use a quadratic function, since we would always get a second-degree polynomial in the original coordinates
- ▶ We use others that we briefly review here


# Activation Functions

## ReLU Function

### ► ReLU Function

- The most popular choice, due to both simplicity of implementation and its good performance on a variety of predictive tasks, is the rectified linear unit (ReLU).
- ReLU provides a very simple nonlinear transformation. Given an element  $x$ , the function is defined as the maximum of that element and 0:

$$\text{ReLU}(x) = \max\{x, 0\}.$$




# Activation Functions

## Sigmoid Function (Logit)

- ▶ The sigmoid function transforms its inputs, for which values lie in the domain  $\mathbb{R}$ , to outputs that lie on the interval  $(0, 1)$ .
- ▶ For that reason, the sigmoid is often called a squashing function: it squashes any input in the range  $(-\infty, \infty)$  to some value in the range  $(0, 1)$ :

$$\text{sigmoid}(x) = \frac{1}{1 + \exp(-x)}.$$



- ▶ In the earliest neural networks, scientists were interested in modeling biological neurons which either fire or do not fire. Thus the pioneers of this field,



# Activation Functions

## ► Other Activation functions

► Tanh:  $\tanh(x) = \frac{1 - \exp(-2x)}{1 + \exp(-2x)}$ .

► Hard tanh:  $\max(-1, \min(1, x))$

► Radial basis function (RBF):  $\exp\left(\frac{1}{\sigma^2} ||W - x||^2\right)$

► Softplus:  $\log(1 + e^x)$

►  $h = \cos(Wx + b)$  Goodfellow et al. (2016) claim that on the MNIST dataset they obtained an error rate of less than 1 percent

► Hidden unit design remains an active area of research, and many useful hidden unit types remain to be discovered

# Output Functions

- ▶ The choice of cost function is tightly coupled with the choice of output unit.
- ▶ Most of the time, we simply use the distance between the data distribution and the model distribution.
  - ▶ Linear  $y = W'h + b \rightarrow \mathbb{R}$
  - ▶ Sigmoid (Logistic)  $\frac{1}{1+\exp(-x)} \rightarrow \text{classification } \{0, 1\}$
  - ▶ Softmax  $\frac{\exp(x)}{\sum \exp(x)} \rightarrow \text{classification multiple categories}$

# Architecture Design

- ▶ Another key design consideration for neural networks is determining the architecture.
- ▶ The word architecture refers to the overall structure of the network: how many units it should have and how these units should be connected to each other.
- ▶ The universal approximation theorem guarantees that regardless of what function we are trying to learn, a sufficiently large MLP will be able to represent this function.

# Architecture Design

- ▶ We are not guaranteed, however, that the training algorithm will be able to learn that function.
- ▶ Even if the network is able to represent the function, learning can fail for two different reasons.
  - 1 The optimization algorithm used for training may not be able to find the value of the parameters that corresponds to the desired function.
  - 2 The training algorithm might choose the wrong function as a result of overfitting

# Architecture Design

- ▶ A feedforward network with a single layer is sufficient to represent any function, but the layer may be infeasibly large and may fail to learn and generalize correctly.
- ▶ In many circumstances, using deeper models can reduce the number of units required to represent the desired function and can reduce the amount of generalization error.
- ▶ The ideal network architecture for a task must be found via experimentation guided by monitoring the validation set error

## Further Readings

- ▶ Aston Zhang, Zachary C. Lipton, Mu Li, and Alexander J. Smola (2020) Dive into Deep Learning. Release 0.15.1. <http://d2l.ai/index.html>
- ▶ Goodfellow, I., Bengio, Y., Courville, A., & Bengio, Y. (2016). Deep learning (Vol. 1, No. 2). Cambridge: MIT press. <http://www.deeplearningbook.org>
- ▶ James, G., Witten, D., Hastie, T., & Tibshirani, R. (2021). An introduction to statistical learning.
- ▶ Rstudio (2020). Tutorial TensorFlow [https://tensorflow.rstudio.com/tutorials/beginners/basic-ml/tutorial\\_basic\\_classification/](https://tensorflow.rstudio.com/tutorials/beginners/basic-ml/tutorial_basic_classification/)
- ▶ Taddy, M. (2019). Business data science: Combining machine learning and economics to optimize, automate, and accelerate business decisions. McGraw Hill Professional.